

This lab is to deal with **Logistic Regression**, **kNN**, and **Decision Tree** algorithms applied to classification tasks.

- **Deadline: 23:59, 01/04/2024**

✓ Import libraries

```
# code
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.impute import SimpleImputer
from sklearn import preprocessing
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression, LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn import metrics
from sklearn.metrics import accuracy_score, classification_report
```

✓ Task 1.

Apply **LogisticRegression** to iris dataset to classify species of iris based on sepal_length (chiều dài đài hoa), sepal_width, petal_length (chiều dài cánh hoa), petal_width. The species are '**setosa**', '**versicolor**' and '**virginica**'.

```
from sklearn import datasets
data1 = datasets.load_iris()

# code
from sklearn import datasets
from sklearn.linear_model import LogisticRegression #import linear
from sklearn.metrics import confusion_matrix #import matrixs để đánh giá model
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score

data1 = datasets.load_iris()
X1 = data1.data
y1 = data1.target
#chia tệp thành train và test , test 30%
Xtrain, Xtest, ytrain, ytest = train_test_split(X1, y1, test_size=0.3)
#train
regression = LogisticRegression(random_state = 0)
regression.fit(Xtrain, ytrain)
#predict y
y_pred = regression.predict(Xtest)
print ("Accuracy : ", accuracy_score(ytest, y_pred))
print ("\nClassification Report:")
print(classification_report(ytest, y_pred))

Accuracy :  0.9555555555555556

Classification Report:
              precision    recall  f1-score   support

    0         1.00        1.00        1.00        11
    1         0.94        0.94        0.94        18
    2         0.94        0.94        0.94        16

 accuracy          0.96
 macro avg         0.96
weighted avg         0.96
```

✓ Task 2.

Apply LogisticRegression to **FASHION** dataset (*fashion_train.csv* and *fashion_test.csv*) which aims at classifying 10 fashion categories. Dataset includes 784 pixels values of images (28x28). This pixel-value is an integer between 0 and 255. Each training and test example is assigned to one of the following labels:

- 0 T-shirt/top
- 1 Trouser
- 2 Pullover
- 3 Dress
- 4 Coat
- 5 Sandal
- 6 Shirt
- 7 Sneaker
- 8 Bag
- 9 Ankle boot

```
from google.colab import drive
drive.mount('/content/gdrive')
%cd '/content/gdrive/MyDrive/machine_learning/lab4'
```

```
Mounted at /content/gdrive
/content/gdrive/MyDrive/machine_learning/lab4
```

```
# code
train_data = pd.read_csv('fashion_train.csv')
test_data = pd.read_csv('fashion_test.csv')
X_train1 = train_data.iloc[:, :784]
y_train1 = train_data.iloc[:, -1]
```

```
X_test1 = test_data.iloc[:, :784]
y_test1 = test_data.iloc[:, -1]
# train model sử dụng train set
model1 = LogisticRegression(max_iter=1000)
model1.fit(X_train1, y_train1)
y_pred1 = model1.predict(X_test1)
```

```
#đánh giá mô hình
accuracy1 = accuracy_score(y_test1, y_pred1)
print ("Accuracy : ", accuracy1)
print("\nClassification Report:")
print(classification_report(y_test1, y_pred1))
```

```
Accuracy :  0.783
```

```
Classification Report:
              precision    recall  f1-score   support

     0           0.72       0.82      0.77         91
     1           0.95       0.96      0.95         92
     2           0.59       0.71      0.65         91
     3           0.88       0.77      0.82        105
     4           0.66       0.66      0.66         99
     5           0.88       0.78      0.83        105
     6           0.51       0.43      0.47         99
     7           0.83       0.88      0.86         94
     8           0.93       0.93      0.93        115
     9           0.86       0.86      0.86        109

 accuracy          0.78         0.78      0.78       1000
  macro avg       0.78         0.78      0.78       1000
 weighted avg     0.79         0.78      0.78       1000
```

✓ Task 3.

Apply another classification algorithm named **kNN**, which is an instance classification model.

- 3.1. Perform kNN algorithm to Iris dataset with $k=\{3, 5, \dots, 29\}$. Select the best value of k . Plot the values of **accuracy, precision, recall, f1 measure** metrics with different values of k .
- 3.2. Then compare the obtained results with those using Logistic regression (based on metrics: **accuracy, precision, recall, f1 measure**) using **PrettyTable**.

```
# task 3.1
import numpy as np
from sklearn import datasets
from sklearn.neighbors import KNeighborsClassifier
from sklearn import neighbors
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

data1 = datasets.load_iris()
X3 = data1.data
y3= data1.target

#tách data , 30% test 70% train
X_train2, X_test2, y_train2, y_test2 = train_test_split(
    X3, y3, test_size=30)

#kiểm tra với k range từ 3->29 step 2
k_values = range(3, 30, 2)
# tệp lưu đánh giá
accuracies=[]
precisions=[]
recalls=[]
f1_scores=[]
for k in k_values:
    knn_model = KNeighborsClassifier(n_neighbors=k)
    # train với tệp 50 50 ở trên
    knn_model.fit(X_train2, y_train2)
    # predic y
    y_pred2 = knn_model.predict(X_test2)
    # đánh giá
    accuracy = accuracy_score(y_test2, y_pred2)
    precision = precision_score(y_test2, y_pred2, average='macro')
    recall = recall_score(y_test2, y_pred2, average='macro')
    f1 = f1_score(y_test2, y_pred2, average='macro')
    #thêm vào tệp lưu
    accuracies.append(accuracy)
    precisions.append(precision)
    recalls.append(recall)
    f1_scores.append(f1)

#lấy giá trị tốt nhất
best_accuracy = np.argmax(accuracies)
best_k = k_values[best_accuracy]


print('giá trị tốt nhất k:',best_k ,'- với accuracy',accuracies[best_k])
# show với k
plt.figure(figsize=(10, 6))
plt.plot(k_values, accuracies, label='Accuracy')
plt.plot(k_values, precisions, label='Precision')
plt.plot(k_values, recalls, label='Recall')
plt.plot(k_values, f1_scores, label='F1 Score')
plt.xlabel('k')
plt.ylabel('Score')
plt.xticks(k_values)
plt.legend()
plt.grid(True)
plt.show()

#3.2 so sánh k best
from prettytable import PrettyTable

t = PrettyTable(['', 'KNeighbors', 'Logistic regression'])
t.add_row(['Accuracy', accuracies[best_k], accuracy_score(ytest, y_pred)])
t.add_row(['precision', precisions[best_k], metrics.precision_score(ytest, y_pred, average='macro')])
t.add_row(['F1', f1_scores[best_k], metrics.f1_score(ytest, y_pred, average='macro')])
```

```
t.add_row(['recall', recalls[best_k], metrics.recall_score(ytest, y_pred, average='macro')])
print(t)
```

giá trị tốt nhất k: 3 - với accuracy 0.9333333333333333



	KNeighbors	Logistic regression
Accuracy	0.9333333333333333	0.9555555555555556
precision	0.9388888888888888	0.9606481481481483
F1	0.9388888888888888	0.9606481481481483
recall	0.9388888888888888	0.9606481481481483

Task 4.

Similar to Task 3, apply kNN algorithm to **FASHION** dataset which included in datasets of sklearn API.

- 4.1. Perform kNN algorithm to Iris dataset with $k=\{3, 5, \dots, 29\}$. Select the best value of k . Plot the values of **accuracy, precision, recall, f1 measure** metrics with different values of k .
- 4.2. Plot the values of **accuracy, precision, recall, f1 measure** metrics with different values of k .
- 4.3. Then compare the obtained results with those using Logistic regression (based on metrics: accuracy, precision, recall, f1 measure).

```
# code
train_data4 = pd.read_csv('fashion_train.csv')
test_data4 = pd.read_csv('fashion_test.csv')
X_train4 = train_data4.iloc[:, :784]
y_train4 = train_data4.iloc[:, -1]

X_test4 = test_data4.iloc[:, :784]
y_test4 = test_data4.iloc[:, -1]

#kiểm tra với k range từ 3->29 step 2
k4_values = range(3, 30, 2)
# tệp lưu đánh giá
accuracies4 = []
precisions4 = []
recalls4 = []
f1_scores4 = []
for a in k4_values:
    knn4_model = KNeighborsClassifier(n_neighbors=a)
    # train với tệp 50 50 ở trên
    knn4_model.fit(X_train4, y_train4)
    # predic y
    y_pred4 = knn4_model.predict(X_test4)
    # đánh giá
```

```


accuracy = accuracy_score(y_test4, y_pred4)
precision = precision_score(y_test4, y_pred4, average='macro')
recall = recall_score(y_test4, y_pred4, average='macro')
f1 = f1_score(y_test4, y_pred4, average='macro')
#thêm vào tệp lưu
accuracies4.append(accuracy)
precisions4.append(precision)
recalls4.append(recall)
f1_scores4.append(f1)

#lấy giá trị tốt nhất
best_accuracy4 = np.argmax(accuracies4)
best_k4 = k4_values[best_accuracy4]

print('giá trị tốt nhất k:', best_k4, '- với accuracy', accuracies4[best_k4])
#k=5 tốt nhất
# show với k
plt.figure(figsize=(10, 6))
plt.plot(k4_values, accuracies4, label='Accuracy')
plt.plot(k4_values, precisions4, label='Precision')
plt.plot(k4_values, recalls4, label='Recall')
plt.plot(k4_values, f1_scores4, label='F1 Score')
plt.xlabel('k')
plt.ylabel('Score')
plt.xticks(k4_values)
plt.legend()
plt.grid(True)
plt.show()
# vẽ pretty so sánh với regression
tb = PrettyTable(['', 'KNeighbors', 'Logistic regression'])
tb.add_row(['Accuracy', accuracies4[5], accuracy_score(y_test1, y_pred1)])
tb.add_row(['precision', precisions4[5], metrics.precision_score(y_test1, y_pred1, average='macro')])
tb.add_row(['F1', f1_scores4[5], metrics.f1_score(y_test1, y_pred1, average='macro')])
tb.add_row(['recall', recalls4[5], metrics.recall_score(y_test1, y_pred1, average='macro')])
print(tb)

```

📄 giá trị tốt nhất k: 5 - với accuracy 0.752



	KNeighbors	Logistic regression
Accuracy	0.752	0.783
precision	0.768657281971095	0.7805722822135391
F1	0.7470186833739232	0.7790413311115769
recall	0.7536326207068555	0.7814072147995884

Task 5.

Compare the performance of selected classification algorithms (**Decision Tree, kNN, and Logistic Regression**) to **spam detection**. The dataset can be accessed from the link: <http://archive.ics.uci.edu/ml/datasets/Spambase> Attribute Information: The last column of 'spambase.csv' denotes whether the e-mail was considered **spam (1) or not (0)**, i.e. unsolicited commercial e-mail. Most of the attributes indicate whether a particular word or character was frequently occurring in the e-mail. The run-length attributes (55-57) measure the length of sequences of consecutive capital letters. For the statistical measures of each attribute, see the end of this file. Here are the definitions of the attributes:

- **48 continuous real [0,100] attributes** of type word_freq_WORD = percentage of words in the e-mail that match WORD, i.e. $100 * (\text{number of times the WORD appears in the e-mail}) / \text{total number of words in e-mail}$. A "word" in this case is any string of alphanumeric characters bounded by non-alphanumeric characters or end-of-string. **Example:** word_freq_address: percentage of words in the e-mail that match ADDRESS.
- **6 continuous real [0,100] attributes** of type char_freq_CHAR = percentage of characters in the e-mail that match CHAR, i.e. $100 * (\text{number of CHAR occurrences}) / \text{total characters in e-mail}$
- **1 continuous real [1,...] attribute** of type capital_run_length_average = average length of uninterrupted sequences of capital letters
- **1 continuous integer [1,...] attribute** of type capital_run_length_longest = length of longest uninterrupted sequence of capital letters
- **1 continuous integer [1,...] attribute** of type capital_run_length_total = sum of length of uninterrupted sequences of capital letters = total number of capital letters in the e-mail
- **1 nominal {0,1} class attribute** of type spam = denotes whether the e-mail was considered spam (1) or not (0), i.e. unsolicited commercial e-mail.

In order to compare the performance of selected algorithms, some common metrics including **accuracy, precision, recall, f1 measures** could be used.

```
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
from sklearn.tree import DecisionTreeClassifier
```

```
# code
```

```
data5 = pd.read_csv('spambase.csv')
X5 = data5.iloc[:, :-1]
y5 = data5.iloc[:, -1]
```

```
#chia data 30-70
```

```
X_train5, X_test5, y_train5, y_test5 = train_test_split(X5, y5, test_size=0.3)
```

```
# train logistic
```

```
logistic_model5 = LogisticRegression(random_state=0);
logistic_model5.fit(X_train5, y_train5)
y_logis_pred5 = logistic_model5.predict(X_test5)
#đánh giá
logis_accuracy = accuracy_score(y_test5, y_logis_pred5)
logis_precision = precision_score(y_test5, y_logis_pred5, average='macro')
logis_recall = recall_score(y_test5, y_logis_pred5, average='macro')
logis_f1 = f1_score(y_test5, y_logis_pred5, average='macro')
```

```
#train decision
```

```
decision_model5 = DecisionTreeClassifier(random_state=42)
decision_model5.fit(X_train5, y_train5)
y_deci_pred5 = decision_model5.predict(X_test5)
#đánh giá
deci_accuracy = accuracy_score(y_test5, y_deci_pred5)
deci_precision = precision_score(y_test5, y_deci_pred5, average='macro')
deci_recall = recall_score(y_test5, y_deci_pred5, average='macro')
deci_f1 = f1_score(y_test5, y_deci_pred5, average='macro')
```

```
#train knn
```

```
knn_model5 = KNeighborsClassifier()
knn_model5.fit(X_train5, y_train5)
y_knn_pred5 = knn_model5.predict(X_test5)
#đánh giá
knn_accuracy = accuracy_score(y_test5, y_knn_pred5)
knn_precision = precision_score(y_test5, y_knn_pred5, average='macro')
knn_recall = recall_score(y_test5, y_knn_pred5, average='macro')
knn_f1 = f1_score(y_test5, y_knn_pred5, average='macro')
```

```
from prettytable import PrettyTable
```

```
# These 3 are the columns of the tables
```

```
t = PrettyTable(["", 'decision', 'Knn', 'logistic'])
```

```
# To insert rows:
```

```
t.add_row(['Accuracy', deci_accuracy, knn_accuracy, logis_accuracy])
t.add_row(['Precision', deci_precision, knn_precision, logis_precision])
t.add_row(['Recall', deci_recall, knn_recall, logis_recall])
t.add_row(['f1 score', deci_f1, knn_f1, logis_f1])
```

```
l.add_row([ f1_score , deci_tf, knn_tf, logis_tf])
print(t)
```

	decision	Knn	logistic
Accuracy	0.9022447501810282	0.8073859522085446	0.9225199131064447
Precision	0.895798721129566	0.7960547122074637	0.9187182910547396
Recall	0.8982628649978937	0.7988759062590072	0.917205063964703
f1_score	0.8969868224245876	0.7973724041762644	0.9179475646695932

Finally,

Save a copy in your Github. Remember renaming the notebook.